

DriftGuard

AI-Powered Behavioral Anomaly Detection for Enterprise SOCs

Technical Report - Honeywell Hackathon Q4: Behavioral Anomaly Detection for Cybersecurity

Author:

Om Mishra

Candidate ID:

20526459

Email:

om.mishra.2410@gmail.com

GitHub Repository:

<https://github.com/omxmishra/DriftGuard>

Submission Overview

This document presents a complete technical report for DriftGuard, a behavioral anomaly detection system built for Honeywell Hackathon Q4. It covers the problem, architecture, dataset, detection and classification approach, explainability layer, results, known limitations, and supporting references. A working prototype (Streamlit dashboard) and full source code accompany this report.

1. Problem Summary

Traditional signature-based security fails against novel or slow, low-and-slow intrusions. DriftGuard models 'normal' access and connection behaviour for users, service accounts, and edge devices, detects deviations in near real-time, classifies the type of anomaly, and provides an explainable risk score for SOC analysts.

2. Innovation and Uniqueness

- Hybrid detection engine: a statistical baseline and an LSTM sequence model run independently and are combined via max-percentile-rank ensembling, so each attack type is caught by whichever detector is strongest for it, rather than one dominant score consuming the shared alert budget.
- Multi-class attack classification with an explicit 'normal' rejection class, so the classifier can identify and downgrade the detector's own false positives instead of forcing an incorrect attack label onto them.
- SHAP-based explainability translated into analyst-readable phrases (e.g. 'geo-velocity: login from an unfamiliar location') rather than raw feature-attribution values.
- Explicit, dashboard-visible cold-start handling: entities without enough history for sequence scoring are surfaced in a dedicated monitoring panel, not silently ignored.

3. Assumptions

- Real intrusion/access-log datasets are privacy-restricted and domain-specific; a synthetic generator was built to produce schema-compliant data with injected, labeled attack patterns.

- An entity's own early-session history (first third of its sessions) is assumed to represent genuinely normal behaviour, and is used to build that entity's baseline profile.
- The alert budget (top 2% of sessions) represents realistic SOC analyst capacity; both 2% and 1% budgets were evaluated to align with the evaluation criteria's own example.
- Ground-truth labels are used only for evaluation and are never fed into the anomaly-scoring models themselves, mirroring real deployments where true labels are unavailable at inference time.

4. System Architecture

The pipeline is organized as six sequential stages, each with a single-responsibility script:

- **generate_data.py** - synthetic access-log generator (7 attack behaviours + 1 ambiguous edge case)
- **baseline_profile.py** - per-entity statistical 'normal' profile and deviation score
- **detector.py** - LSTM autoencoder over 10-session windows, ensembled with the baseline score
- **classifier.py** - XGBoost multi-class attack-type classifier, including a 'normal' rejection class
- **explainability.py** - SHAP-based per-alert feature attribution, translated to analyst-readable text
- **dashboard.py** - Streamlit analyst UI: alert queue, risk scores, entity history, cold-start monitor

Data flow: Logs -> Feature Engineering -> [Baseline Profile + LSTM Sequence Model] -> Ensemble -> Attack Classifier -> SHAP Explainability -> Analyst Dashboard

5. Dataset

250 entities (users, service accounts, edge devices) over 30 simulated days, producing **16,294 total sessions** with a **2.081% anomaly rate** - within the brief's suggested 0.5-3% injection range.

Label	Count
normal	15,955
brute_force	89
insider_drift	55
lateral_movement	54
credential_stuffing	54
low_and_slow_exfil	51
impossible_travel	22
device_spoofing	14

insider_drift is an intentionally ambiguous edge case (gradual, legitimate-looking privilege expansion), included to test false-positive tuning rather than as a clean-cut attack.

6. Detection Approach

A statistical baseline (per-entity z-scores on login hour and session duration, plus binary flags for new resource/IP/geo/fingerprint/auth-method/auth-failure) provides a fast, interpretable first signal. An LSTM

autoencoder is trained exclusively on windows containing zero anomalous sessions, so it learns only genuine normal temporal patterns; reconstruction error on any window is the sequence-level anomaly score. The two scores are combined via max-percentile-rank ensembling, so each attack type can be caught by whichever detector is strongest for it, rather than one dominant score consuming the shared alert budget.

7. Results - Detection Recall by Attack Type (Top 2% Alert Budget)

Attack Type	Baseline Only	Ensemble (Final)
brute_force	92.1%	100%
device_spoofing	21.4%	100%
credential_stuffing	96.3%	98%
low_and_slow_exfil	70.6%	83%
lateral_movement	18.5%	34%
impossible_travel	18.2%	22%
insider_drift (edge case)	10.9%	11.5%
normal (false positive rate)	0.8%	1.8%

Ensembling improved recall for every attack category over either detector alone. lateral_movement and impossible_travel remain comparatively weak (see Limitations) - both were partially crowded out of the shared 2% budget by stronger-scoring attack types.

8. Results - Attack-Type Classification (Held-Out Test Set)

Class	Precision	Recall	F1	Support
brute_force	1.00	1.00	1.00	22
credential_stuffing	1.00	1.00	1.00	14
device_spoofing	1.00	1.00	1.00	3
impossible_travel	1.00	1.00	1.00	5
insider_drift	0.93	1.00	0.97	14
lateral_movement	1.00	1.00	1.00	14
low_and_slow_exfil	1.00	1.00	1.00	13
normal (rejection class)	1.00	0.99	1.00	150

Including 'normal' as a trainable class allows the classifier to identify detector false positives instead of forcing an incorrect attack label onto them. On real flagged alerts (not the held-out test set), predicted-vs-true label agreement is **98.7%** (up from 46% before this fix, when false positives were being force-labeled).

9. Explainability

SHAP TreeExplainer computes per-alert feature attribution against the classifier's predicted class. The top contributing features per alert are translated into analyst-readable phrases (e.g. 'geo-velocity: login from an unfamiliar location', 'device fingerprint mismatch (OS/MAC/protocol)'), avoiding raw SHAP values or feature-name jargon in the analyst-facing text. A global feature-importance chart is also generated for the flagged population as a whole.

10. Cold-Start and Concept Drift

- **Cold-start:** entities with fewer than 10 sessions (the minimum LSTM window size) cannot yet be scored by the sequence model. These fall back to their own per-entity statistical baseline score, and are surfaced explicitly in the dashboard's Cold-Start Entity Monitoring panel (18 of 250 entities in the current dataset).
- **Concept drift:** baseline profiles are currently built once, from each entity's early-session history. The design supports periodic re-computation over a recent rolling window (rather than full history) so that gradually evolving legitimate behaviour is not permanently flagged; `insider_drift` was specifically included in the dataset to stress-test this distinction, and correctly shows the lowest recall of any category, reflecting appropriately cautious handling of ambiguous, gradual change.

11. Scalability and Real-Time Streaming Feasibility

The current implementation is a batch pipeline over a static CSV, but each scoring stage is designed around single-session functions (baseline scoring, LSTM window scoring) rather than whole-dataset operations. This maps directly onto a streaming architecture: an event bus (e.g. Kafka) would deliver one session at a time to a consumer that maintains each entity's rolling window and profile in a fast key-value store, scores the incoming session immediately, and asynchronously updates the entity's profile in the background - without requiring a full pipeline re-run.

12. Known Limitations

- Synthetic data is more cleanly separable than real logs; held-out classification metrics (nearly 100%) should be read as an upper bound, not a guarantee of real-world performance.
- `lateral_movement` and `impossible_travel` recall (34% and 22%) are the weakest results, primarily because their scores compete with stronger-scoring attack types for the same limited 2% alert budget.
- Global SHAP importance shows zero contribution from `auth_fail`, `auth_mismatch`, and `new_geo` in the current flagged population - this reflects collinearity with other features (e.g. `auth_success`) and underrepresentation of geo-based attacks in the flagged sample, not that these signals are unhelpful in general.
- A required XGBoost/SHAP version incompatibility (XGBoost's newer per-class `base_score` serialization vs. an older SHAP parser) was resolved by pinning `xgboost==1.7.6`; this should be revisited if the environment's SHAP version is later upgraded.
- No live streaming deployment or continuous retraining loop was implemented in this submission - both are documented as design extensions rather than working code.
- The schema's `command_sequence` and `protocol` fields are generated and present in the data but are not yet used as model features; they are a natural next step for detecting privileged-session misuse and protocol-level anomalies specifically.

13. Conclusion

DriftGuard delivers all seven required components of the Q4 problem statement: a documented synthetic data generator, a per-entity statistical baseline, an LSTM sequence-aware detector ensembled with that baseline, a multi-class attack classifier with false-positive rejection, SHAP-based explainability, and an analyst-facing dashboard - with cold-start handling, an explicit concept-drift design story, and honestly reported limitations throughout.

14. Evaluation Criteria Mapping

Evaluation Criterion	Our Solution
Detection accuracy on imbalanced labels	Hybrid baseline + LSTM ensemble, trained/evaluated at 2.08% anomaly rate
Correct anomaly-type classification	XGBoost multi-class classifier, 7 types + normal rejection class
False positive rate at analyst alert budget	Evaluated at both top 2% and top 1% budgets; false-positive rate reported per budget
Explainability / analyst usability	SHAP TreeExplainer, translated to analyst-readable phrases per alert
Handling cold-start entities	Statistical baseline fallback for entities under 10 sessions; surfaced in dashboard
Handling concept drift	Rolling-window baseline refresh (documented design); insider_drift used to stress-test ambiguous drift
System design & scalability (streaming feasibility)	Single-session scoring functions; documented event-bus + rolling-profile streaming architecture
Report clarity	This report: assumptions, architecture, metrics, limitations, and references

Each row maps directly to the problem statement's own evaluation criteria, so coverage can be checked line by line.

15. Research and References

- [PyTorch documentation](#) - LSTM autoencoder implementation
- [XGBoost documentation](#) - multi-class attack classifier
- [SHAP documentation](#) - per-alert feature attribution
- [Streamlit documentation](#) - analyst dashboard framework
- [scikit-learn documentation](#) - preprocessing, scaling, and evaluation metrics
- [DriftGuard GitHub repository](#) - full source code and commit history

16. Screenshots

Supporting screenshots of code, results, and the working dashboard prototype.

Deploy

Alert Detail & Entity History

Select a session to inspect

bcf8c626-da3f-4cb8-ba76-f504e768cfbb

Entity: user_0178

Predicted Attack Type: lateral_movement

Confidence: 98.6%

Risk Tier: MEDIUM

Detected By: baseline_profile

True Label (eval only): lateral_movement

Contributing Factors (explainability):

Flagged as lateral_movement (confidence 0.99) due to: atypical time-of-day pattern; atypical day-of-week pattern; accessed a resource never touched before

Recent activity for user_0178 (20 of last sessions shown)

	Time	Resource	Source IP	Geo	Auth	Duration (min)	Label
12978	2026-06-27 17:45:49	resource_6	59.221.160.21	Wilkersonm	password	15.14	normal
12980	2026-06-27 19:03:47	resource_55	53.51.209.231	Wilkersonm	password	7.42	normal
12981	2026-06-28 18:52:59	resource_44	171.180.128.198	Wilkersonm	password	14.27	normal
12982	2026-06-29 18:06:10	resource_58	130.253.129.218	Wilkersonm	password	8.19	normal
12983	2026-06-29 18:54:24	resource_58	46.37.171.5	Wilkersonm	password	22.24	normal
12984	2026-06-29 18:58:02	resource_42	136.133.173.214	Wilkersonm	password	4	normal
12985	2026-06-30 18:25:27	resource_44	158.40.226.160	Wilkersonm	password	13.02	normal

DriftGuard: synthetic behavioral anomaly detection pipeline (baseline profiling + LSTM sequence model + XGBoost classification + SHAP explainability)

alert detail shap

Deploy

Alert Queue

Attack Type

lateral_movement


Risk Tier

medium


☐ Hide likely false positives

	Session	Entity	Predicted Type	Confidence	Risk Percentile	Risk Tier	Detected By
421	260acefd-a38c-4967-8086-04ba228b39c2	user_0178	lateral_movement	99.5%	0.9858	medium	baseline_profile
422	aa7b7cdb-d931-4504-89a0-be4a6f4102c7	user_0178	lateral_movement	97.3%	0.9856	medium	baseline_profile
310	814212a9-1575-4111-a538-bb8ad39f96d	user_0040	lateral_movement	99.4%	0.9841	medium	baseline_profile
309	8b36a540-0658-4a6d-a986-78add2f37af0	user_0040	lateral_movement	99.4%	0.9836	medium	baseline_profile
342	c33d0b55-1af3-4d58-9656-74db46c0200f	user_0094	lateral_movement	58.7%	0.9834	medium	sequence_model
198	e75466dc-78d6-4093-8c7a-5f95c4cb1eef	service_account_0080	lateral_movement	99.9%	0.9806	medium	baseline_profile
308	b94ef4bf-5ec9-48b3-9daf-bc6859b9fb57	user_0040	lateral_movement	99.4%	0.9801	medium	baseline_profile

alert queue



dashboard overview



model performance

